




# Concurrent Programming

Effective Programming in Scala

# Distributed Computations Are Ubiquitous

Nowadays, distributed systems power a lot of applications (e.g., web and mobile applications).

These systems are made of several physical machines (a.k.a. nodes) communicating together:




# Challenges Raised by Distributed Systems

- ▶ What happens if multiple clients modify the state of the system at the same time?
- ▶ Can the server handle requests coming from other clients while it is already busy responding to one client?

## Distributed Computations Are Ubiquitous (2)


If we zoom on the server node, the same pattern might be used internally:



Similar concerns arise: can the system handle the application of two commands at the same time? Is it possible to read the system state while commands are still being processed?

## Distributed Computations Are Ubiquitous (3)

Last, if we zoom again on one node it can use the same pattern again to leverage the multiple CPUs of the machine:



Again: what happens if multiple computations modify the memory at the same time? How can a computation running in a CPU depend on the result of another computation running in a different CPU?

# Multi-Threading

- ▶ Leveraging multiple CPUs in a single program requires **multi-threading**,
- ▶ sharing data between threads requires using **thread-safe** data structures (similarly, sharing data between nodes requires some form of transactions in the underlying database system).

# Sharing Data Between Concurrent Programs

- ▶ Within a single JVM, use thread-safe data structures such as the ones in the package `java.util.concurrent.atomic`, or `scala.collection.concurrent`,
- ▶ Between processes or nodes, you can rely on third-party database systems.

# Dependencies Between Computations

Whether you perform a remote call, or you start a computation in a different thread, you eventually need a way to “do something” with the result.

In Scala, in both situations we use the type Future to model the result of such computations.

A Future represents a value that may not be available yet, but might become available in the future—when whatever is computing it has finished.



## Future

As a first example, consider the following signature of a program that inserts a new user in a remote database and eventually returns a result of type `User`:

```
def insertUser(login: String, passwordHash: Seq[Byte]): Future[User]
```

As a second example, consider the following signature of a program that computes the cryptographic hash of a password using the BCrypt algorithm. This program starts the computation on a different thread and eventually returns a result of type `Seq[Byte]`:

```
def bcrypt(saltRound: Int, password: String): Future[Seq[Byte]]
```

# Asynchronous Computations

We will see later how to implement an operation that executes code on a different execution thread (locally, or remotely). Let us first get a general idea of what it is like to work with operations returning Future.

The operations `insertUser` and `bcrypt` return a result to the caller thread possibly *before* the actual result has been computed.


We say that `bcrypt` and `insertUser` are **asynchronous computations**.

The returned Future value provides methods to eventually use the actual result *when* it is available.

## Manipulating Future Values

The most common way to work with Future values is to **transform** them into other Future values by using operations very similar to the ones we already use with collections or Try: `map`, `flatMap` and `zip`.


For instance, consider the complete scenario of adding a new user to the system. We first want to compute the user password hash with the operation `bcrypt`, and, after we get the password hash, insert the user row in the database with the `insertUser` operation:



The arrow indicates the direction in which the time flows.

## Future Values All the Way Down

By combining Future values, we end up building trees of computations with sequential and parallel branches:



The result is necessarily a Future value again, which is finally consumed by your application framework or testing framework.

# Failure Handling

When we run an asynchronous computation, there are two possible outcomes:

- ▶ the computation eventually returns a result,
- ▶ or, it fails (e.g., an exception is thrown, or the remote service is unreachable).

We will see in the next lesson how to handle such failures.

We can model the state of a `Future[A]` value at any point in time with the type `Option[Try[A]]`:

- ▶ `None` means that the outcome of the computation is not known yet,
- ▶ `Some(Success(a))` means that the computation succeeded,
- ▶ `Some(Failure(e))` means that the computation failed.

# Summary

Most systems run on several execution threads.

Sharing data between several threads of execution requires using thread-safe data structures.

For a program running in a given thread, we model the result of computations possibly performed in a different thread with the type Future.

We usually combine values of type Future to build a graph of computations containing sequential and parallel branches.